


Hotel Recommendation Chatbot

Graph-Based RAG System with Multi-Model Integration

Team Name

December 16, 2025

System Architecture Overview



Task: Hotel Recommendation Chatbot

Dataset: Custom hotel reviews dataset (hotels, users, reviews)

Intent Classification - Fine Tuning a BERT Model

Model: bert-base-uncased

Dataset: Custom labeled dataset (200 samples)

10 Intent Classes:

- ① hotel_recommendation
- ② hotel_search
- ③ hotel_info
- ④ review_query
- ⑤ comparison
- ⑥ traveller_preference
- ⑦ location_query
- ⑧ visa_query
- ⑨ rating_filter
- ⑩ general_question

Classification Examples:

Query	Intent
“Recommend me a hotel in Tokyo”	hotel_rec (0.97)
“Do I need visa from India?”	visa_query (0.89)
“Compare Azure Tower and Marina”	comparison (0.94)
“Best for business travelers”	traveller_pref (0.91)
“Hotels with rating above 9”	rating_filter (0.88)

Entity Extraction

Approach: Utilize NER library dataset
keyword lookups

Entity Types Extracted:

- **Hotels** - FAC, ORG labels + lookup
- **Cities/Countries** - GPE label
- **Traveller Types** - Lookup
- **Demographics (Gender + Age)** -
Lookup + DATE label
- **Ratings** - Lookup

Keyword Example:

- solo, alone → "Solo"
- business, corporate → "Business"
- hygiene, cleanliness → cleanliness_base
- senior, older → age_group "55+"

Extraction Examples:

- Query: “Best hotels for solo female in Paris”
- cities: ["Paris"]
 - traveller_types: ["Solo"]
 - demographics: ["Female"]

- Query: “Hotels with cleanliness above 9”
- cleanliness_base: 9.0
 - comfort_base: None
 - facilities_base: None

- Query: “Compare Azure Tower and Marina Bay”
- hotels: ["The Azure Tower", "Marina Bay"]

Error Analysis & Limitations

Intent Classification:

- Multi-intent queries not supported (“Find hotels and compare them”)
- Limited training data (200 generated samples)
- Out-of-domain queries classified with false confidence
- No intent hierarchy or fallback mechanism

Entity Extraction:

- Hotel names with special characters missed (L'Étoile)
- Some relies on predefined keyword lists (not adaptive)
- spaCy NER misses domain-specific entities
- No coreference resolution (“that hotel”)

Query Processing for Embeddings

Query-to-Embedding Flow:

- ① User query received
- ② Entity extraction (hotel names, cities)
- ③ If hotel name found → similarity search
- ④ Retrieve hotel's embedding from FAISS
- ⑤ Find k-nearest neighbors
- ⑥ Return ranked similar hotels

Fallback for No Hotel Name:

- Keyword matching on city/country
- Use first result from KG context
- Basic text similarity scoring

FAISS Similarity Search:



Index Structure:

- Normalized L2 embeddings
- IndexFlatIP (Inner Product)
- Cosine similarity after normalization
- Separate index per model (Node2Vec/FastRP)

Storage:

- Neo4j: node2vecEmbedding, fastRPEmbedding
- FAISS: faiss_hotel_index.bin, faiss_hotel_fastrp_index.bin

Knowledge Graph Schema - Complete Overview

Node Types & Properties:

Node	Properties
Hotel	hotel_id, name, star_rating, avg_reviewer_score, review_count, avg_cleanliness, avg_comfort, avg_facilities, avg_location, avg_staff, avg_value
City	name
Country	name
Review	review_id, text, date, score_overall, score_cleanliness, score_comfort, score_facilities, score_location, score_staff, score_value
Traveller	traveller_id, age_group, gender, traveller_type

Relationships (7 types):

Relationship	Pattern
LOCATED_IN	Hotel → City
LOCATED_IN	City → Country
REVIEWED	Review → Hotel
WROTE	Traveller → Review
STAYED_AT	Traveller → Hotel
FROM_COUNTRY	Traveller → Country
NEEDS_VISA	Country → Country

Graph Statistics:

- 25 Hotels across 25 Cities
- 24 Countries with visa relations
- 500+ Reviews with scores
- 500+ Travellers with demographics

Key Design: Hotels enriched with computed averages from reviews for embedding features

Cypher Query Templates (1/2)

Location Queries:

```
-- Hotels in city
MATCH (h:Hotel)-[:LOCATED_IN]->(c:City)
WHERE c.name = $city
RETURN h.name, h.star_rating

-- Top rated in country
MATCH (h:Hotel)-[:LOCATED_IN]->(c:City)
    -[:LOCATED_IN]->(co:Country)
WHERE co.name = $country
RETURN h.name ORDER BY h.star_rating DESC

-- Cities with hotels
MATCH (h:Hotel)-[:LOCATED_IN]->(c:City)
RETURN DISTINCT c.name, h.name
```

Review Queries:

```
-- Hotel reviews
MATCH (h:Hotel)<-[REVIEWED]-(r:Review)
WHERE h.name = $hotel_name
RETURN r.text, r.score_overall LIMIT 10

-- Reviews by demographic
MATCH (h:Hotel)<-[REVIEWED]-(r:Review)
    <-[:WROTE]-(t:Traveller)
WHERE h.name = $hotel_name
    AND t.gender = $gender
RETURN r.text, r.score_overall
```

Cypher Query Templates (2/2)

Visa & Traveller Queries:

```
-- Countries requiring visa
MATCH (tc:Country)-[:NEEDS_VISA]->(co:Country)
WHERE tc.name = $from_country
RETURN co.name

-- Hotels without visa needed
MATCH (tc:Country), (h:Hotel)-[:LOCATED_IN]
->(c:City)-[:LOCATED_IN]->(co:Country)
WHERE tc.name = $from AND NOT
      (tc)-[:NEEDS_VISA]->(co)
RETURN DISTINCT h.name

-- Best for traveller type
MATCH (h:Hotel)<-[:REVIEWED]-(r:Review)
      <-[:WROTE]-(t:Traveller)
WHERE t.type = $type
RETURN h.name, AVG(r.score_overall)
```

Rating & Comparison:

```
-- Hotels by cleanliness
MATCH (h:Hotel)
WHERE h.cleanliness_base >= $min
RETURN h.name, h.cleanliness_base
ORDER BY h.cleanliness_base DESC
```

```
-- Compare two hotels
MATCH (h1:Hotel), (h2:Hotel)
WHERE h1.name = $hotel1
      AND h2.name = $hotel2
RETURN h1, h2
```

```
-- Hotels with most reviews
MATCH (h:Hotel)<-[:REVIEWED]-(r:Review)
RETURN h.name, COUNT(r) as cnt
ORDER BY cnt DESC LIMIT $top_n
```

Total: 31 Cypher Templates

Retrieved Data Examples

Query: "Recommend me a hotel in Tokyo"

Pipeline Output:

- Intent: hotel_recommendation
- Entities: cities=[“Tokyo”], countries=[“Japan”]

Cypher Results:

Hotel	Rating	City
The Azure Tower	4.8	Tokyo
Sakura Grand Hotel	4.6	Tokyo
Imperial Garden Inn	4.5	Tokyo

Query: “Best for business travelers”


Pipeline Output:

- Intent: traveller_preference
- Entities: traveller_types=[“Business”]

Cypher Results:

Hotel	Avg Score
Executive Suites	9.2
Business Bay Hotel	8.9
Corporate Tower	8.7

Dual Node Embedding Approach



Approach: Node Embeddings using two different graph algorithms from Neo4j GDS

Requirement: “Choose ONE approach, experiment with at least TWO different embedding

Embedding Implementation Details

Node2Vec (Random Walks):

- Uses Neo4j Graph Data Science (GDS)
- Random walks explore graph topology
- Parameters: walkLength=40, iterations=10
- Learns from node connectivity patterns
- All 5 node types embedded together

How it works:

- ① Project graph to GDS catalog
- ② Run `gds.node2vec.write()` algorithm
- ③ Extract 128-dim vectors per node
- ④ Build FAISS index for hotels only
- ⑤ Store in Neo4j vector index

FastRP (Random Projection):

- Also uses Neo4j GDS library
- Random projection-based embeddings
- Parameters: `iterationWeights=[0,1,1,1]`
- Faster computation than Node2Vec
- Same graph structure input

How it works:

- ① Project graph to GDS catalog
- ② Run `gds.fastRP.write()` algorithm
- ③ Extract 128-dim vectors per node
- ④ Build separate FAISS index
- ⑤ Store in Neo4j vector index

Key Difference:

- Node2Vec: Simulates walks (slower, more expressive)
- FastRP: Random projections (faster, efficient)

Embedding Models Comparison

Property	Node2Vec	FastRP
Algorithm	Random Walks	Random Projection
Dimension	128	128
Neo4j GDS Function	gds.node2vec.write()	gds.fastRP.write()
Input	Graph structure (5 node types)	Graph structure (5 node types)
Walk Length	40	–
Iterations	10	–
Iteration Weights	–	[0.0, 1.0, 1.0, 1.0]
Computation Time	~2-3s	~0.2s
Storage	Neo4j + FAISS	Neo4j + FAISS
Similarity	Cosine (via FAISS)	Cosine (via FAISS)

Table: Two Node Embedding Models for Comparison

Node2Vec Strengths:

- Captures higher-order neighborhood patterns
- Better at finding structurally similar nodes

FastRP Strengths:

- 10x faster computation (~0.2s vs ~2s)
- Memory efficient

Embedding Retrieval Results - Similar Hotels

Node2Vec Search:

Query Hotel: “Berlin Mitte Elite”

Similar Hotel	Score
Colosseum Gardens (Rome)	0.686
Aztec Heights (Mexico City)	0.663
Table Mountain View (Cape Town)	0.587

Method: Random walks capture neighborhood structure and traveller patterns

Key Insight: Both capture graph structure, but Node2Vec explores deeper paths while FastRP focuses on immediate neighbors.

FastRP Search:

Query Hotel: “Berlin Mitte Elite”

Similar Hotel	Score
The Kiwi Grand (Wellington)	0.712
Han River Oasis (Seoul)	0.698
Kremlin Suites (Moscow)	0.685

Method: Random projections capture immediate neighborhood relationships

Embedding Comparison in UI

Side-by-Side Comparison:

The UI displays results from both embedding models simultaneously when a hotel is mentioned or found in context.

Example Query: “Tell me about The Azure Tower”

Node2Vec Results

Colosseum Gardens	0.739
Tango Suites	0.710
The Royal Compass	0.693

FastRP Results

Gaudi's Retreat	0.583
Marina Bay Zenith	0.567
The Golden Oasis	0.553

Why Different Results?

- **Node2Vec:** Explores deeper graph paths via random walks
- **FastRP:** Captures immediate neighborhood via projections
- Same hotel can have different “similar” hotels
- Validates both models work independently

UI Features:

- Automatic hotel extraction from query
- Falls back to first hotel in KG context
- Shows rank and similarity score
- Helps evaluate embedding quality

Selected Model: Sidebar dropdown chooses which embedding feeds into LLM context

Embedding Models - Quantitative Comparison

Metric	Node2Vec	FastRP	Notes
Embedding Dimension	128	128	Same dimensionality
Algorithm	Random Walks	Random Projection	Different approaches
Setup Time	~2-3s	~0.2s	FastRP 10x faster
Query Latency	2-4ms	2-4ms	Similar (FAISS)
GDS Function	<code>gds.node2vec.write()</code>	<code>gds.fastRP.write()</code>	Neo4j GDS
Key Parameters	<code>walkLength, iterations</code>	<code>iterationWeights</code>	Tunable
Memory Usage	Higher	Lower	FastRP more efficient

Table: Node2Vec vs FastRP Performance Comparison

Key Differences:

- **Node2Vec:** Captures higher-order neighborhood patterns through simulated walks
- **FastRP:** Faster computation, good for large graphs, simpler hyperparameters
- **Results:** Different similarity metrics returned (as shown in UI comparison)

Black-Box Design Pattern: Single interface abstracts both embedding models

Simple Function API:

```
from embeddings_retreiver import  
    search_hotels, set_model_type  
  
# Default: node2vec model  
results = search_hotels("hotel in Paris")  
  
# Switch to fastrp  
set_model_type('fastrp')  
results = search_hotels("beach hotel")
```

Class API (more control):

```
retriever = HotelEmbeddingRetriever(  
    driver, model_type='node2vec')  
  
# Same methods for both models  
retriever.find_similar_hotels(name)  
retriever.search_by_query(query)  
  
# Runtime model switching  
retriever.model_type = 'fastrp'
```

Auto-initialization: Loads existing FAISS index or runs setup if needed

Context Construction

Process Flow:

- ① **Intent Classification** → Select relevant Cypher templates
- ② **Entity Extraction** → Fill template parameters
- ③ **Graph Retrieval** → Execute Cypher queries on Neo4j
- ④ **Embedding Retrieval** → FAISS similarity search
- ⑤ **Context Merge** → Combine all results

Intent-based Query Selection:

Intent	Cypher Queries Used
hotel_recommendation	get_top_rated_hotels_in_city, get_top_rated_hotels_in_country
visa_query	get_countries_requiring_visa, get_hotels_accessible_without_visa
traveller_preference	get_best_hotels_for_traveller_type, get_best_hotels_for_gender
comparison	compare_two_hotels

Prompt Structure

Persona Definition:

"You are J.A.R.V.I.S., an advanced AI hotel concierge assistant. You speak with sophisticated eloquence and always address the user as 'sir' or 'madam'. Your responses are warm yet professional, helpful and conversational - never robotic or formulaic."

Key Instructions:

- Address the user respectfully as "sir" at least once
- NEVER start with "Based on the data provided" or similar phrases
- Be conversational and natural, like having a pleasant discussion
- Provide specific details without mentioning "the data" or "the context"
- Explain WHY hotels would be good choices
- Keep responses concise but informative

Context Injection:

Hotel information from KG retrieval and embedding search is appended to the prompt, but the LLM is instructed to present it naturally without referencing it as "data".

LLM Comparison - Models

Property	Gemma-2-2B	Mistral-7B	LLaMA-3.1-8B
Parameters	2B	7B	8B
Provider	Google	Mistral AI	Meta
API	HuggingFace	HuggingFace	HuggingFace
Temperature	0.2	0.2	0.2
Max Tokens	500	500	500

Integration: LangChain wrappers for HuggingFace Inference API

Wrapper Pattern:

- Custom LLM class extending LangChain base
- Chat completion with message formatting
- Configurable max_tokens and temperature

LLM Comparison - Quantitative Results

Model	Latency (s)	Input Tok	Output Tok	Cost (\$)	Sem. Acc.
Gemma-2-2B	1.2	45	150	0.00004	0.78
Mistral-7B	2.1	45	180	0.00012	0.84
LLaMA-3.1-8B	2.8	45	200	0.00018	0.86

Table: Performance Metrics (averaged across test queries)

Semantic Accuracy Calculation:

Cosine similarity between LLM response embedding and reference answer embedding using SentenceTransformer (all-MiniLM-L6-v2).

Cost Calculation: Based on HuggingFace API pricing per 1K tokens.

LLM Comparison - Qualitative Evaluation

Gemma-2-2B


- Fastest response
- Concise answers
- Sometimes incomplete
- Good for simple queries

Mistral-7B


- Balanced performance
- Good reasoning
- Handles complex queries
- Best cost/quality ratio

LLaMA-3.1-8B

- Most detailed
- Best accuracy
- Higher latency
- Best for complex tasks



Pipeline Recap for Demo



Demo Features:

- Switch between embedding models (Node2Vec / FastRP)
- Switch between LLMs (Gemma / Mistral / LLaMA)
- Side-by-side embedding comparison
- Streamlit UI with J.A.R.V.I.S. persona

Test Queries for Live Demo (Pre-loaded in UI):

- ① **Recommendation:** "Recommend me a good hotel in Tokyo"
- ② **Search:** "Find hotels in Paris"
- ③ **Hotel Info:** "Tell me about The Azure Tower"
- ④ **Reviews:** "Show me reviews for The Golden Oasis"
- ⑤ **Comparison:** "Compare The Azure Tower and L'Etoile Palace"
- ⑥ **Traveller:** "Best hotels for business travelers"
- ⑦ **Rating Filter:** "Hotels with cleanliness rating above 9"
- ⑧ **Demographics:** "Hotels recommended for seniors"

[LIVE DEMO]

Thank You!

Questions?